

```
In [139]: import pandas as pd
import numpy as np
import re
from sklearn.preprocessing import normalize
from sklearn.feature_extraction.text import TfidfVectorizer
from tqdm import tqdm
import matplotlib.pyplot as plt
import os
from scipy.sparse import hstack
import warnings
import seaborn as sns
```

```
In [102]: dfnlp = pd.read_csv('NLP_features_train.csv',encoding='latin-1')
dfppro = pd.read_csv('df_fe_without_preprocessing_train.csv',encoding='latin-1')
```

```
In [103]: df1 = dfnlp.drop(['qid1', 'qid2'],axis=1)
df2 = dfppro.drop(['qid1', 'qid2', 'question1', 'question2', 'is_duplicate'],axis=1)
```

```
In [104]: df1.head()
```

```
Out[104]:
```

	id	question1	question2	is_duplicate	cwc_min	cwc_max	csc_min	csc_max	ctc_min	ctc_max
0	0	what is the step by step guide to invest in sh...	what is the step by step guide to invest in sh...	0	0.999980	0.833319	0.999983	0.999983	0.916667	0.785714
1	1	what is the story of kohinoor koh i noor dia...	what would happen if the indian government sto...	0	0.799984	0.399996	0.749981	0.599988	0.700000	0.466667
2	2	how can i increase the speed of my internet co...	how can internet speed be increased by hacking...	0	0.399992	0.333328	0.399992	0.249997	0.400000	0.285714
3	3	why am i mentally very lonely how can i solve...	find the remainder when math 23 24 math i...	0	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000
4	4	which one dissolve in water quickly sugar salt...	which fish would survive in salt water	0	0.399992	0.199998	0.999950	0.666644	0.571429	0.307692

```
In [105]: df2.head()
```

```
Out[105]:
```

	id	freq_qid1	freq_qid2	q1len	q2len	q1_n_words	q2_n_words	word_common	word_Total	W
0	0	1	1	66	57	14	12	10.0	23.0	
1	1	4	1	51	88	8	13	4.0	20.0	
2	2	1	1	73	59	14	10	4.0	24.0	
3	3	1	1	50	65	11	9	0.0	19.0	
4	4	3	1	76	39	13	7	2.0	20.0	


```
In [106]: df1 = df1.merge(df2,on='id',how='left')
```

```
In [107]: df1.head(2)
```

```
Out[107]:
```

	id	question1	question2	is_duplicate	cwc_min	cwc_max	csc_min	csc_max	ctc_min	ctc_max
0	0	what is the step by step guide to invest in sh...	what is the step by step guide to invest in sh...	0	0.999980	0.833319	0.999983	0.999983	0.916667	0.785714
1	1	what is the story of kohinoor koh i noor dia...	what would happen if the indian government sto...	0	0.799984	0.399996	0.749981	0.599988	0.700000	0.460526

2 rows × 30 columns



```
In [108]: df1 = df1[df1['question1'].notnull()]
df1 = df1[df1['question2'].notnull()]
```

```
In [110]: Y = df1['is_duplicate']
Y.shape
```

```
Out[110]: (404269,)
```

```
In [112]: df1 = df1.drop(['id', 'is_duplicate'], axis=1)
```

```
In [113]: tfidfq1 = TfidfVectorizer()
q1_tfidf = tfidfq1.fit_transform(df1['question1'].values.astype('U'))
```

```
In [114]: type(q1_tfidf)
```

```
Out[114]: scipy.sparse.csr.csr_matrix
```

```
In [115]: q1_tfidf.shape
```

```
Out[115]: (404269, 67909)
```

```
In [116]: tfidfq2 = TfidfVectorizer()
q2_tfidf = tfidfq2.fit_transform(df1['question2'].values.astype('U'))
```

```
In [117]: q2_tfidf.shape
```

```
Out[117]: (404269, 62704)
```

```
In [118]: ## Combining the two questions together
q_tfidf = hstack((q1_tfidf, q2_tfidf))
```

```
In [119]: df1 = df1.drop(['question1', 'question2'], axis=1, inplace=True)
```

```
In [120]: q_tfidf.shape
```

```
Out[120]: (404269, 130613)
```

```
In [121]: type(q_tfidf)
```

```
Out[121]: scipy.sparse.coo.coo_matrix
```

```
In [155]: from sklearn.metrics import confusion_matrix
from sklearn.metrics.classification import accuracy_score, log_loss
from sklearn.feature_extraction.text import TfidfVectorizer
from collections import Counter
from scipy.sparse import hstack
from sklearn.multiclass import OneVsRestClassifier
from sklearn.svm import SVC
from sklearn.model_selection import StratifiedKFold
from collections import Counter, defaultdict
from sklearn.calibration import CalibratedClassifierCV
from sklearn.naive_bayes import MultinomialNB
from sklearn.naive_bayes import GaussianNB
from sklearn.model_selection import train_test_split
from sklearn.model_selection import RandomizedSearchCV
import math
from sklearn.metrics import normalized_mutual_info_score
from sklearn.ensemble import RandomForestClassifier

from sklearn.model_selection import cross_val_score
from sklearn.linear_model import SGDClassifier
##from mlxtend.classifier import StackingClassifier

from sklearn import model_selection
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import precision_recall_curve, auc, roc_curve
```

```
In [156]: X_train, X_test, Y_train, Y_test = train_test_split(q_tfidf, Y, stratify=Y, test_size=0
```

```
In [124]: print(X_train.shape)
print(X_test.shape)
```

```
(282988, 130613)
```

```
(121281, 130613)
```

```

In [125]: def plot_confusion_matrix(test_y, predict_y):
    C = confusion_matrix(test_y, predict_y)
    # C = 9,9 matrix, each cell (i,j) represents number of points of class i are

    A = (((C.T)/(C.sum(axis=1))).T)
    #divid each element of the confusion matrix with the sum of elements in that

    # C = [[1, 2],
    #      [3, 4]]
    # C.T = [[1, 3],
    #        [2, 4]]
    # C.sum(axis = 1) axis=0 corresonds to columns and axis=1 corresponds to row
    # C.sum(axix =1) = [[3, 7]]
    # ((C.T)/(C.sum(axis=1))) = [[1/3, 3/7]
    #                             [2/3, 4/7]]

    # ((C.T)/(C.sum(axis=1))).T = [[1/3, 2/3]
    #                               [3/7, 4/7]]
    # sum of row elements = 1

    B = (C/C.sum(axis=0))
    #divid each element of the confusion matrix with the sum of elements in that
    # C = [[1, 2],
    #      [3, 4]]
    # C.sum(axis = 0) axis=0 corresonds to columns and axis=1 corresponds to row
    # C.sum(axix =0) = [[4, 6]]
    # (C/C.sum(axis=0)) = [[1/4, 2/6],
    #                       [3/4, 4/6]]
    plt.figure(figsize=(20,4))

    labels = [1,2]
    # representing A in heatmap format
    cmap=sns.light_palette("blue")
    plt.subplot(1, 3, 1)
    sns.heatmap(C, annot=True, cmap=cmap, fmt=".3f", xticklabels=labels, yticklab
    plt.xlabel('Predicted Class')
    plt.ylabel('Original Class')
    plt.title("Confusion matrix")

    plt.subplot(1, 3, 2)
    sns.heatmap(B, annot=True, cmap=cmap, fmt=".3f", xticklabels=labels, yticklab
    plt.xlabel('Predicted Class')
    plt.ylabel('Original Class')
    plt.title("Precision matrix")

    plt.subplot(1, 3, 3)
    # representing B in heatmap format
    sns.heatmap(A, annot=True, cmap=cmap, fmt=".3f", xticklabels=labels, yticklab
    plt.xlabel('Predicted Class')
    plt.ylabel('Original Class')
    plt.title("Recall matrix")

    plt.show()

```

## Logistic Regression(Hyper-parameter tuned)

```
In [131]: alpha = [10**x for x in range(-5,2)]

log_error_array=[]
for i in alpha:
    clf = SGDClassifier(alpha=i,penalty='l1',loss='hinge',random_state=42)
    clf.fit(X_train,Y_train)
    sig_clf = CalibratedClassifierCV(clf,method="sigmoid")
    sig_clf.fit(X_train,Y_train)
    predict_y = sig_clf.predict_proba(X_test)
    log_error_array.append(log_loss(Y_test,predict_y,labels=clf.classes_,eps=1e-1)
    print('For the values of alpha = ',i,"The log loss is:",log_loss(Y_test,predi
```

D:\Anaconda\envs\tensorflow\lib\site-packages\sklearn\linear\_model\stochastic\_gradient.py:128: FutureWarning: max\_iter and tol parameters have been added in <class 'sklearn.linear\_model.stochastic\_gradient.SGDClassifier'> in 0.19. If both are left unset, they default to max\_iter=5 and tol=None. If tol is not None, max\_iter defaults to max\_iter=1000. From 0.21, default max\_iter will be 1000, and default tol will be 1e-3.


"and default tol will be 1e-3." % type(self), FutureWarning)

D:\Anaconda\envs\tensorflow\lib\site-packages\sklearn\linear\_model\stochastic\_gradient.py:128: FutureWarning: max\_iter and tol parameters have been added in <class 'sklearn.linear\_model.stochastic\_gradient.SGDClassifier'> in 0.19. If both are left unset, they default to max\_iter=5 and tol=None. If tol is not None, max\_iter defaults to max\_iter=1000. From 0.21, default max\_iter will be 1000, and default tol will be 1e-3.

"and default tol will be 1e-3." % type(self), FutureWarning)

D:\Anaconda\envs\tensorflow\lib\site-packages\sklearn\linear\_model\stochastic\_gradient.py:128: FutureWarning: max\_iter and tol parameters have been added in <class 'sklearn.linear\_model.stochastic\_gradient.SGDClassifier'> in 0.19. If both are left unset, they default to max\_iter=5 and tol=None. If tol is not None, max\_iter defaults to max\_iter=1000. From 0.21, default max\_iter will be 1000, and default tol will be 1e-3.

```
In [133]: fig,ax = plt.subplots()
ax.plot(alpha,log_error_array,c='g')
for i,txt in enumerate(np.round(log_error_array,3)):
    ax.annotate((alpha[i],np.round(txt,3)),(alpha[i],log_error_array[i]))
plt.grid()
plt.title("Cross Validation Error for each alpha")
plt.xlabel("Alpha i's")
plt.ylabel("Error measure")
plt.show()
```



```
In [135]: best_alpha = np.argmin(log_error_array)
clf = SGDClassifier(alpha=alpha[best_alpha],penalty='l2',loss='log',random_state=
clf.fit(X_train,Y_train)
sig_clf = CalibratedClassifierCV(clf,method="sigmoid")
sig_clf.fit(X_train,Y_train)
```

D:\Anaconda\envs\tensorflow\lib\site-packages\sklearn\linear\_model\stochastic\_gradient.py:128: FutureWarning: max\_iter and tol parameters have been added in < class 'sklearn.linear\_model.stochastic\_gradient.SGDClassifier'> in 0.19. If both are left unset, they default to max\_iter=5 and tol=None. If tol is not None, max\_iter defaults to max\_iter=1000. From 0.21, default max\_iter will be 1000, and default tol will be 1e-3.

"and default tol will be 1e-3." % type(self), FutureWarning)

D:\Anaconda\envs\tensorflow\lib\site-packages\sklearn\linear\_model\stochastic\_gradient.py:128: FutureWarning: max\_iter and tol parameters have been added in < class 'sklearn.linear\_model.stochastic\_gradient.SGDClassifier'> in 0.19. If both are left unset, they default to max\_iter=5 and tol=None. If tol is not None, max\_iter defaults to max\_iter=1000. From 0.21, default max\_iter will be 1000, and default tol will be 1e-3.

"and default tol will be 1e-3." % type(self), FutureWarning)

D:\Anaconda\envs\tensorflow\lib\site-packages\sklearn\linear\_model\stochastic\_gradient.py:128: FutureWarning: max\_iter and tol parameters have been added in < class 'sklearn.linear\_model.stochastic\_gradient.SGDClassifier'> in 0.19. If both are left unset, they default to max\_iter=5 and tol=None. If tol is not None, max\_iter defaults to max\_iter=1000. From 0.21, default max\_iter will be 1000, and default tol will be 1e-3.

"and default tol will be 1e-3." % type(self), FutureWarning)

D:\Anaconda\envs\tensorflow\lib\site-packages\sklearn\linear\_model\stochastic\_gradient.py:128: FutureWarning: max\_iter and tol parameters have been added in < class 'sklearn.linear\_model.stochastic\_gradient.SGDClassifier'> in 0.19. If both are left unset, they default to max\_iter=5 and tol=None. If tol is not None, max\_iter defaults to max\_iter=1000. From 0.21, default max\_iter will be 1000, and default tol will be 1e-3.

"and default tol will be 1e-3." % type(self), FutureWarning)

```
Out[135]: CalibratedClassifierCV(base_estimator=SGDClassifier(alpha=1e-05, average=False,
class_weight=None, epsilon=0.1,
eta0=0.0, fit_intercept=True, l1_ratio=0.15,
learning_rate='optimal', loss='log', max_iter=None, n_iter=None,
n_jobs=1, penalty='l2', power_t=0.5, random_state=42, shuffle=True,
tol=None, verbose=0, warm_start=False),
cv=3, method='sigmoid')
```




```
In [140]: predict_y = sig_clf.predict_proba(X_train)
print('For values of best alpha = ', alpha[best_alpha], "The train log loss is:",
predict_y = sig_clf.predict_proba(X_test)
print('For values of best alpha = ', alpha[best_alpha], "The test log loss is:",1
predicted_y =np.argmax(predict_y,axis=1)
print("Total number of data points :", len(predicted_y))
plot_confusion_matrix(Y_test, predicted_y)
```

For values of best alpha = 1e-05 The train log loss is: 0.4720810857646023

For values of best alpha = 1e-05 The test log loss is: 0.5082508140455909

Total number of data points : 121281



## Linear SVM with hyperparameter tuning

```
In [142]: alpha = [10 ** x for x in range(-5, 2)]
log_error_array=[]
for i in alpha:
    clf = SGDClassifier(alpha=i, penalty='l1', loss='hinge', random_state=42)
    clf.fit(X_train, Y_train)
    sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
    sig_clf.fit(X_train, Y_train)
    predict_y = sig_clf.predict_proba(X_test)
    log_error_array.append(log_loss(Y_test, predict_y, labels=clf.classes_, eps=1e-15))
print('For values of alpha = ', i, "The log loss is:", log_loss(Y_test, predict_y, labels=clf.classes_, eps=1e-15))
```

D:\Anaconda\envs\tensorflow\lib\site-packages\sklearn\linear\_model\stochastic\_gradient.py:128: FutureWarning: max\_iter and tol parameters have been added in <class 'sklearn.linear\_model.stochastic\_gradient.SGDClassifier'> in 0.19. If both are left unset, they default to max\_iter=5 and tol=None. If tol is not None, max\_iter defaults to max\_iter=1000. From 0.21, default max\_iter will be 1000, and default tol will be 1e-3.


"and default tol will be 1e-3." % type(self), FutureWarning)

D:\Anaconda\envs\tensorflow\lib\site-packages\sklearn\linear\_model\stochastic\_gradient.py:128: FutureWarning: max\_iter and tol parameters have been added in <class 'sklearn.linear\_model.stochastic\_gradient.SGDClassifier'> in 0.19. If both are left unset, they default to max\_iter=5 and tol=None. If tol is not None, max\_iter defaults to max\_iter=1000. From 0.21, default max\_iter will be 1000, and default tol will be 1e-3.

"and default tol will be 1e-3." % type(self), FutureWarning)

D:\Anaconda\envs\tensorflow\lib\site-packages\sklearn\linear\_model\stochastic\_gradient.py:128: FutureWarning: max\_iter and tol parameters have been added in <class 'sklearn.linear\_model.stochastic\_gradient.SGDClassifier'> in 0.19. If both are left unset, they default to max\_iter=5 and tol=None. If tol is not None, max\_iter defaults to max\_iter=1000. From 0.21, default max\_iter will be 1000, and default tol will be 1e-3.

```
In [143]: fig, ax = plt.subplots()
ax.plot(alpha, log_error_array, c='g')
for i, txt in enumerate(np.round(log_error_array, 3)):
    ax.annotate((alpha[i], np.round(txt, 3)), (alpha[i], log_error_array[i]))
plt.grid()
plt.title("Cross Validation Error for each alpha")
plt.xlabel("Alpha i's")
plt.ylabel("Error measure")
plt.show()
```



```
In [145]: best_alpha = np.argmin(log_error_array)
clf = SGDClassifier(alpha=alpha[best_alpha], penalty='l1', loss='hinge', random_s
clf.fit(X_train, Y_train)
sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
sig_clf.fit(X_train, Y_train)
```

D:\Anaconda\envs\tensorflow\lib\site-packages\sklearn\linear\_model\stochastic\_gradient.py:128: FutureWarning: max\_iter and tol parameters have been added in < class 'sklearn.linear\_model.stochastic\_gradient.SGDClassifier'> in 0.19. If both are left unset, they default to max\_iter=5 and tol=None. If tol is not None, max\_iter defaults to max\_iter=1000. From 0.21, default max\_iter will be 1000, and default tol will be 1e-3.

"and default tol will be 1e-3." % type(self), FutureWarning)

D:\Anaconda\envs\tensorflow\lib\site-packages\sklearn\linear\_model\stochastic\_gradient.py:128: FutureWarning: max\_iter and tol parameters have been added in < class 'sklearn.linear\_model.stochastic\_gradient.SGDClassifier'> in 0.19. If both are left unset, they default to max\_iter=5 and tol=None. If tol is not None, max\_iter defaults to max\_iter=1000. From 0.21, default max\_iter will be 1000, and default tol will be 1e-3.

"and default tol will be 1e-3." % type(self), FutureWarning)

D:\Anaconda\envs\tensorflow\lib\site-packages\sklearn\linear\_model\stochastic\_gradient.py:128: FutureWarning: max\_iter and tol parameters have been added in < class 'sklearn.linear\_model.stochastic\_gradient.SGDClassifier'> in 0.19. If both are left unset, they default to max\_iter=5 and tol=None. If tol is not None, max\_iter defaults to max\_iter=1000. From 0.21, default max\_iter will be 1000, and default tol will be 1e-3.

"and default tol will be 1e-3." % type(self), FutureWarning)


D:\Anaconda\envs\tensorflow\lib\site-packages\sklearn\linear\_model\stochastic\_gradient.py:128: FutureWarning: max\_iter and tol parameters have been added in < class 'sklearn.linear\_model.stochastic\_gradient.SGDClassifier'> in 0.19. If both are left unset, they default to max\_iter=5 and tol=None. If tol is not None, max\_iter defaults to max\_iter=1000. From 0.21, default max\_iter will be 1000, and default tol will be 1e-3.

"and default tol will be 1e-3." % type(self), FutureWarning)

```
Out[145]: CalibratedClassifierCV(base_estimator=SGDClassifier(alpha=1e-05, average=False,
class_weight=None, epsilon=0.1,
eta0=0.0, fit_intercept=True, l1_ratio=0.15,
learning_rate='optimal', loss='hinge', max_iter=None, n_iter=None,
n_jobs=1, penalty='l1', power_t=0.5, random_state=42, shuffle=True,
tol=None, verbose=0, warm_start=False),
cv=3, method='sigmoid')
```

```
In [146]: predict_y = sig_clf.predict_proba(X_train)
print('For values of best alpha = ', alpha[best_alpha], "The train log loss is:",
predict_y = sig_clf.predict_proba(X_test)
print('For values of best alpha = ', alpha[best_alpha], "The test log loss is:",
predicted_y = np.argmax(predict_y,axis=1)
print("Total number of data points :", len(predicted_y))
plot_confusion_matrix(Y_test, predicted_y)
```

For values of best alpha = 1e-05 The train log loss is: 0.49949532720743306  
 For values of best alpha = 1e-05 The test log loss is: 0.521583251161899  
 Total number of data points : 121281



## XGBOOST

```
In [159]: import xgboost as xgb
from xgboost import XGBClassifier
params = {
    'max_depth':[3,5,6,7,8],
    'learning_rate':[0.01,0.02,0.03,0.1,0.2,0.3],
    'n_estimators':[100,200,300,400,500],
    'gamma':[0,0.5,1,1.5,2,5]
}
model = XGBClassifier(nthread=-1)
kfold = StratifiedKFold(n_splits=4,shuffle=True)
random_search = RandomizedSearchCV(model,param_distributions=params,scoring="neg_
random_result = random_search.fit(X_train,Y_train)
```

```
In [160]: print(random_search.best_estimator_)
```

```
XGBClassifier(base_score=0.5, booster='gbtree', colsample_bylevel=1,
              colsample_bytree=1, gamma=2, learning_rate=0.3, max_delta_step=0,
              max_depth=6, min_child_weight=1, missing=None, n_estimators=200,
              n_jobs=1, nthread=-1, objective='binary:logistic', random_state=0,
              reg_alpha=0, reg_lambda=1, scale_pos_weight=1, seed=None,
              silent=True, subsample=1)
```

```
In [161]: print(random_search.best_params_)
```

```
{'n_estimators': 200, 'learning_rate': 0.3, 'max_depth': 6, 'gamma': 2}
```

```
In [162]: import xgboost as xgb
params = {}
params['objective'] = 'binary:logistic'
params['eval_metric'] = 'logloss'
params['n_estimators'] = 200
params['learning_rate'] = 0.3
params['max_depth'] = 6
params['gamma'] = 2

d_train = xgb.DMatrix(X_train,label=Y_train)
d_test = xgb.DMatrix(X_test,label=Y_test)

watch_list = [(d_train,'train'),(d_test,'valid')]
bst = xgb.train(params,d_train,400,watch_list,early_stopping_rounds=20,verbose_eval=True)
xgdmatrix = xgb.DMatrix(X_train,Y_train)
```

```
[0] train-logloss:0.652804 valid-logloss:0.653483
Multiple eval metrics have been passed: 'valid-logloss' will be used for early
stopping.
```

Will train until valid-logloss hasn't improved in 20 rounds.

```
[10] train-logloss:0.564021 valid-logloss:0.56849
[20] train-logloss:0.540598 valid-logloss:0.547044
[30] train-logloss:0.526073 valid-logloss:0.534579
[40] train-logloss:0.515753 valid-logloss:0.525502
[50] train-logloss:0.507798 valid-logloss:0.519003
[60] train-logloss:0.500988 valid-logloss:0.513297
[70] train-logloss:0.494719 valid-logloss:0.508467
[80] train-logloss:0.490062 valid-logloss:0.505168
[90] train-logloss:0.484666 valid-logloss:0.50088
[100] train-logloss:0.47955 valid-logloss:0.497495
[110] train-logloss:0.475843 valid-logloss:0.494796
[120] train-logloss:0.47244 valid-logloss:0.492377
[130] train-logloss:0.468747 valid-logloss:0.489726
[140] train-logloss:0.465923 valid-logloss:0.487911
[150] train-logloss:0.463325 valid-logloss:0.486273
[160] train-logloss:0.460787 valid-logloss:0.484735
[170] train-logloss:0.457638 valid-logloss:0.482706
[180] train-logloss:0.455027 valid-logloss:0.481064
[190] train-logloss:0.451385 valid-logloss:0.478542
[200] train-logloss:0.449522 valid-logloss:0.477311
[210] train-logloss:0.44727 valid-logloss:0.475891
[220] train-logloss:0.445042 valid-logloss:0.474544
[230] train-logloss:0.442694 valid-logloss:0.472868
[240] train-logloss:0.440596 valid-logloss:0.471629
[250] train-logloss:0.438466 valid-logloss:0.470357
[260] train-logloss:0.434884 valid-logloss:0.467995
[270] train-logloss:0.431886 valid-logloss:0.465711
[280] train-logloss:0.430291 valid-logloss:0.464759
[290] train-logloss:0.428458 valid-logloss:0.46379
[300] train-logloss:0.427097 valid-logloss:0.463067
[310] train-logloss:0.42553 valid-logloss:0.462268
[320] train-logloss:0.423604 valid-logloss:0.461015
[330] train-logloss:0.420737 valid-logloss:0.459352
[340] train-logloss:0.419403 valid-logloss:0.458584
```


```
[350] train-logloss:0.41737 valid-logloss:0.457503
[360] train-logloss:0.415963 valid-logloss:0.456825
[370] train-logloss:0.414507 valid-logloss:0.456152
[380] train-logloss:0.413223 valid-logloss:0.455474
[390] train-logloss:0.411768 valid-logloss:0.454846
[399] train-logloss:0.410378 valid-logloss:0.453996
```

```
In [163]: predict_y1 = bst.predict(d_test)
print("The test log Loss is:", log_loss(Y_test, predict_y1, labels=clf.classes_, eps=
```

The test log Loss is: 0.45399573660014325

```
In [164]: predicted_y = np.array(predict_y1>0.5, dtype=int)
print("Total number of data points :", len(predicted_y))
plot_confusion_matrix(Y_test, predicted_y)
```

Total number of data points : 121281



```
In [168]: from prettytable import PrettyTable

x = PrettyTable()
x.field_names = ["Model", "Train_log_loss", "Test_log_loss"]
x.add_row(["Logistic Regression", 0.4720810857646023, 0.5082508140455909])
x.add_row(["Linear SVM", 0.4994953272074330, 0.521583251161899])
x.add_row(["XGBoost", 0.410378, 0.45399573660014325])
print(x)
```

Model	Train_log_loss	Test_log_loss
Logistic Regression	0.4720810857646023	0.5082508140455909
Linear SVM	0.499495327207433	0.521583251161899
XGBoost	0.410378	0.45399573660014325

```
In [ ]:
```